

Spatiotemporal Flood Risk Modeling Under a Single-Orbit Synthetic Aperture Radar Constraint: Confound-Controlled Susceptibility and Flood-Aware Emergency Routing

MUSTAFA BASTHIKODI¹, ([MEMBERSHIP GRADE], IEEE), KRISH P. CHOWTA¹, ADVAITH SHAJEENDRAN¹,
DHRITHI K. S.¹, and GAMAN A. RAI¹

¹Department of Computer Science and Engineering, Sahyadri College of Engineering and Management, Adyar, Mangaluru 575007, Karnataka, India

Corresponding author: Mustafa Basthikodi (e-mail: [INSTITUTIONAL EMAIL]).

This work was supported by [FUNDING SOURCE / grant number, or delete this sentence if the work received no external funding].

ABSTRACT Coastal districts of southwestern India flood every monsoon, yet the satellite record available for mapping those floods is thinner than the modeling literature usually assumes. Over Dakshina Kannada, Karnataka, a single Sentinel-1 orbit samples the district on a nominal twelve-day cycle, so individual flash-flood peaks may occur entirely inside the revisit gap. We therefore decompose risk into a static susceptibility surface, trained on a multi-temporal radar flood-frequency inventory, and a dynamic rainfall trigger. A remote-sensing confound is identified and removed: a model given land cover and a vegetation index reaches an area under the curve of 0.9878, but attribution analysis shows it has learned that closed canopy hides water from radar, not flood physics. The terrain-and-rainfall model scores 0.9625 on a random split, 0.9474 under spatial-block cross-validation, and 0.8594 in a coastal geographic holdout; its highest-susceptibility tenth contains 49.26 percent of observed flood pixels. District-wide operational evaluation uses a 25,407-node, 59,501-edge road graph, 19 curated hospital origins across nine taluks, and 13 scenario points. At the standardized May 2025 setting, 192 of 247 hospital–scenario pairs remain reachable and 55 are severed; among reachable pairs, flood-aware routing reduces mean exposure from 0.1712 to 0.0655 at a mean 63.075-minute detour. A* explores 13.7 percent fewer nodes than Dijkstra but is 21.3 percent slower in this implementation. A district-spread 3×3 dispatch QUBO matches the exact assignment, while QAOA performance is non-monotone with depth and no quantum advantage is claimed.

INDEX TERMS Emergency dispatch, explainable artificial intelligence, flood-aware routing, flood susceptibility, spatial cross-validation, synthetic aperture radar.

I. INTRODUCTION

Flooding is the most frequently recurring natural hazard in India, and the southwestern coastal belt concentrates the mechanisms that make it difficult to forecast at a useful spatial scale. Dakshina Kannada district, centred on Mangaluru in coastal Karnataka, receives more than 4000 mm of rainfall in an average monsoon and drains the Nethravathi and the Gurupura into a shared estuary across a coastal plain narrow enough that fluvial discharge and tidal stage interact within a few kilometres of the built-up area. Extreme-rainfall frequency over this coast is trending upward [1], land-use change in the Nethravathi basin has altered its runoff response [2, 3], and the district's land cover has itself changed measurably [4], so the hazard is neither stationary nor purely meteorological. Two recent events frame this study: the August 2024 Nethravathi river flood and the May 2025 Mangaluru urban flood. Exist-

ing work on this basin is static and index-based, for example the geographic information system (GIS) and analytic-hierarchy zoning of Nirmala [5].

The dominant response in the literature is flood-susceptibility mapping: train a classifier on terrain and hydrological conditioning factors against an inventory of observed flooding [6–9], increasingly paired with explainable-AI attribution so the map can be interrogated rather than only displayed [10–12]. This work sits in that tradition and departs from it in three ways.

The observational premise is usually left implicit: Susceptibility studies source an inventory from whatever observations exist, then treat the resulting map as a model of flooding in general. The density of that record constrains what the map can claim. Over Dakshina Kannada, Sentinel-1 [13] cover-

age is a *single* orbit geometry, a descending pass on relative orbit 63, revisiting approximately every twelve days; we verified this against the Copernicus archive instead of assuming the constellation's global six-day figure applies. The consequence is strict: no individual flood here can be mapped from synthetic aperture radar (SAR), because a coastal flash flood peaks and recedes well inside the revisit gap, whereas the rapid SAR inundation mapping demonstrated elsewhere in India [14] presumes an acquisition near the peak. Section IV-A sets out the decomposition that follows.

Validation is usually optimistic: Random k -fold cross-validation on geospatial samples leaks information across fold boundaries, because neighboring points share terrain, so the score measures interpolation within sampled neighborhoods rather than skill at an unvisited location [15–17]. AUCs above 0.95 are common in this literature and rarely carry a spatial control. We report random k -fold, spatial-block k -fold, and a geographic-transfer test with training entirely outside the evaluation window (Section V-C).

Free parameters are usually asserted: Flood-aware routing introduces a penalty weight and an impassability threshold, typically fixed by judgement and never swept. They determine which roads the router refuses to use, so they are exactly what an operational deployment would be questioned on. Section V-F reports the full 7×7 sweep.

A. CONTRIBUTIONS

1. A two-stage risk decomposition, $R(\mathbf{x}, t) = S(\mathbf{x}) \cdot T(t)$, in which a SAR-flood-frequency-trained static surface is modulated by an independently validated live trigger. The decomposition is derived from a measured satellite-revisit constraint rather than adopted for convenience.
2. Identification and ablation of a remote-sensing circularity that inflates susceptibility scores in forested coastal catchments, where SAR cannot observe water under canopy and a classifier given land cover learns detectability rather than hazard.
3. A three-level validation ladder, random split, spatial-block cross-validation, and geographic transfer, quantifying how much reported skill survives each successive removal of spatial leakage.
4. A complete sensitivity analysis of the routing cost function that distinguishes a genuinely insensitive parameter from one adjacent to a topological discontinuity.
5. A search-transparent Dijkstra/A* benchmark across 192 hospital-scenario pairs, reporting a wall-clock result that runs counter to the node-count result and explaining why.

Fig. 2 maps the full pipeline. A hybrid quadratic unconstrained binary optimization (QUBO) and quantum approximate optimization algorithm (QAOA) dispatch layer is also

implemented and verified against the exact classical optimum. It is reported in Section V-H as a feasibility check rather than as a contribution: at the problem size studied here the classical assignment solver is exact and effectively instantaneous, and no quantum advantage is claimed.

II. RELATED WORK

A. SUSCEPTIBILITY MAPPING AND ITS VALIDATION

Tree ensembles dominate the field. Tehrany et al. [6] established the support vector machine (SVM) and frequency-ratio baseline; Khosravi et al. [7] benchmarked multi-criteria against machine-learning approaches; Arabameri et al. [8] and Saha et al. [9] established random forest and boosted-tree variants as the practical default. Reported AUCs cluster between 0.85 and 0.95. Wang et al. [18] and Avand et al. [19] extended the approach to data-scarce basins. The conditioning factors are largely standardized: elevation, slope, curvature, the topographic wetness index [20], height above nearest drainage [21, 22], distance to channel, drainage density and rainfall climatology.

Explainability entered the field recently. Lundberg and Lee [10] introduced Shapley additive explanations (SHAP); Aydin and Iban [11] and Seydi et al. [12] applied it to flood susceptibility, generally to confirm that elevation and channel proximity dominate. Our SHAP analysis is used differently, as a diagnostic that exposed a confound (Section IV-E), not as a post-hoc confirmation.

That geospatial models are routinely over-validated is established but under-applied. Roberts et al. [15] formalized block cross-validation for spatially structured data; Ploton et al. [16] demonstrated that ignoring it inflates reported accuracy in remote-sensing regressions; Meyer and Pebesma [17] extended the argument to area-of-applicability estimation. Flood-susceptibility papers reporting spatial CV remain a minority, which is part of the motivation for Section V-C.

B. SAR FLOOD MAPPING

SAR change detection for inundation is well established [23–25], typically with backscatter thresholding against a dry reference and a permanent-water mask [26]. The recognized limitation is that C-band cannot detect water beneath closed canopy [27, 28], which is precisely the mechanism behind the confound we isolate. Revisit-limited inventories are usually handled by aggregating across many passes into a frequency product [26, 29], the approach adopted here. Closest to the present work, Sadhwani and Eldho [30] combine SAR with elevation data in a dual machine-learning scheme for inundation mapping and depth estimation over the Indian west-coast belt, using the same sensor and terrain inputs as this study. Their depth-regression stage is precisely the extension our Limitation 1 identifies as the highest-value next step, and the contrast is instructive: they target an absolute physical quantity, while the index reported here is explicitly relative.

Recent event-stratified experiments by Thasma et al. [31]

compare SAR-only, SAR–DEM and SAR–land-cover-prior segmentation across multiple backbones and seeds, reinforcing the need to report event-to-event transfer rather than only pooled accuracy.

C. ROUTING, DISPATCH AND QUANTUM OPTIMIZATION

Risk-weighted shortest-path routing under flooding has an established literature [33, 34]. A* [35] remains the standard goal-directed refinement of Dijkstra [36], with admissibility guaranteeing optimality. Emergency-vehicle siting and allocation is itself a mature operations-research field, from maximal-covering formulations [37] to the ambulance location and relocation models surveyed by Brotcorne et al. [38]; our contribution is not a new formulation but the coupling of that assignment to a SAR-grounded, time-varying risk surface. Ambulance-to-incident assignment is a classical linear assignment problem solved exactly by the Hungarian method [39], which is why any quantum treatment at this scale must be framed as feasibility, never as advantage. Temporal planning for coordinated rescue, medical, evacuation and logistics actions has also been demonstrated on synthetic flood benchmarks by Siddiqui et al. [32]; its stated need for real operational validation motivates the conservative feasibility framing used here.

QAOA [40] approximates combinatorial optima on gate-model hardware; Lucas [41] gives the standard Ising encodings, including assignment. Zhou et al. [42] introduced the INTERP parameter-transfer heuristic used here. Known obstacles, barren plateaus [43] and noise-limited depth, mean that near-term demonstrations at practical scale remain out of reach; our scoping in Section V–H reflects that.

D. GAP

No prior work we are aware of combines an inventory that is explicit about single-orbit revisit limits, a trigger validated against documented events without fitting, a routing layer with a complete parameter sweep and a search-transparent algorithm comparison, and a dispatch layer verified against the exact classical optimum, applied to one Indian coastal district with every number reproducible from public data.

III. STUDY AREA AND DATA

A. STUDY AREA

Dakshina Kannada district occupies the coastal plain and lower Western Ghats scarp of southwestern Karnataka (Fig. 1). The modeling domain is the bounding box [74.776108°, 12.460086°] to [75.670642°, 13.185163°]. Routing is evaluated over the resulting district-wide road graph; geographic transfer uses a coastal SAR footprint within the district. Table I summarizes the physical setting as measured from the assembled data stack, not quoted from secondary sources.

Three features of the setting drive the methodology. First, relief is extreme over short distances: elevation in the sam-

pled floodable domain spans −5 m to 279 m, with the Ghats escarpment rising steeply inland, so runoff concentration times are short and flood peaks are flashy. Second, the coastal strip is low and tidally influenced, so the estuary can impound river discharge, a coupling that a rainfall-only trigger captures only indirectly. Third, land cover in the low-lying non-flooded fraction of the district is predominantly closed-canopy vegetation, which is the mechanism underlying the SAR detectability confound examined in Section IV-E.

B. DOCUMENTED EVENTS

Two events anchor the temporal validation. The August 2024 Nethravathi flood followed sustained catchment wetting, 216.9 mm over the preceding 72 h, with a comparatively modest 53.4 mm in the final 24 h. The May 2025 Mangaluru urban flood combined a high 24-hour total (87.2 mm, IMD “Heavy”) with a similarly saturated antecedent state (205.6 mm/72 h). Their contrasting structure, one saturation-driven, one intensity-driven, is what makes them a useful pair for testing a trigger that weights both. Full inputs are in Table VIII.

C. DATA SOURCES

All inputs are public and ingested programmatically (Table II). Satellite and terrain layers are processed server-side on Google Earth Engine [44]; only sampled and aggregated outputs are downloaded.

IV. METHODOLOGY

A. WHY THE RISK MODEL IS DECOMPOSED

A fully spatiotemporal flood model would estimate inundation $I(\mathbf{x}, t)$ jointly in space and time. Fitting such a model requires observations resolved at the timescale of the process. With one usable SAR orbit at a twelve-day cadence, the observational record cannot constrain the temporal dimension at all. We therefore factor the problem:

$$R(\mathbf{x}, t) = S(\mathbf{x}) \cdot T(t), \quad S(\mathbf{x}) \in [0, 1], \quad T(t) \in [0, 1], \quad (1)$$

where $S(\mathbf{x})$ is a time-invariant relative propensity estimated from the *frequency* of observed flooding across all available passes, and $T(t)$ is a spatially uniform scalar estimated from meteorological data available at hourly cadence. The factorization assigns each term to the data that can actually constrain it. Its cost is an explicit modeling assumption, that the spatial pattern of relative risk does not change with event magnitude, which we return to in Section V-J.

B. SAR FLOOD-FREQUENCY INVENTORY

For each monsoon window (1 June–30 September 2024; 15 May–30 September 2025), every Sentinel-1 GRD scene matching the fixed orbit geometry is speckle-filtered with a

TABLE 1: Study-area characteristics. Terrain and rainfall statistics are computed over the sampled floodable-lowland domain ($n = 3000$ points); network statistics are from the district-wide OpenStreetMap drivable graph.

Property	Value	Source
District	Dakshina Kannada, Karnataka, India	FAO GAUL
Modeling bounding box	74.776108°–75.670642°E, 12.460086°–13.185163°N	OSM Nominatim
Operational raster	120 m, nine factors	this study
Principal rivers	Nethravathi, Gurupura (shared estuary)	MERIT Hydro
Elevation range (sampled)	–5 m to 279 m	SRTM
HAND range (sampled)	0 m to 79.4 m	MERIT Hydro
Mean annual rainfall	3373 mm to 4401 mm (mean 4124)	CHIRPS 2020–2024
Monsoon regime	Southwest, June–September	IMD
Sentinel-1 coverage	1 orbit (descending, rel. orbit 63)	Copernicus
Nominal SAR revisit	12 d	Copernicus
Road graph	25,407 nodes, 59,501 directed edges	OpenStreetMap
Flood-prone edges ($S > 0.5$)	9886 (16.61%)	this study
Hospital routing origins	19 across nine taluks	curated OSM inventory
Scenario points	13	this study
OSM snapshot / audit date	11 September 2026	this study

SRTM: Shuttle Radar Topography Mission; CHIRPS: Climate Hazards group Infrared Precipitation with Stations; IMD: India Meteorological Department; FAO GAUL:




FIGURE 1. Study area. (a) India with Karnataka shaded and Mangaluru marked. (b) Karnataka with the Dakshina Kannada district domain. (c) District-wide hillshaded terrain, channel network, road coverage, 19 curated hospital routing origins, and 13 scenario points. Inclusion as a routing origin does not verify emergency specialty, bed availability, or real-time service capability.

30 m focal median and thresholded for open water:

$$W_k(\mathbf{x}) = \mathbb{1}[\sigma_{VV,k}^0(\mathbf{x}) < -16 \text{ dB}] \cdot \mathbb{1}[\sigma_{VH,k}^0(\mathbf{x}) < -22 \text{ dB}], \quad (2)$$

where $\sigma_{VV,k}^0$ and $\sigma_{VH,k}^0$ are calibrated backscatter coefficients for pass k . A pixel counts as *flooded* on pass k only if it is water then and dry in a January–February baseline, and is not permanent water:

$$F_k(\mathbf{x}) = W_k(\mathbf{x}) \cdot (1 - W_{\text{dry}}(\mathbf{x})) \cdot \mathbb{1}[O_{\text{JRC}}(\mathbf{x}) < 80\%], \quad (3)$$

with O_{JRC} the JRC Global Surface Water occurrence [26]. The inventory is the normalized frequency over the $K = 20$ qualifying passes:

$$\Phi(\mathbf{x}) = \frac{1}{K} \sum_{k=1}^K F_k(\mathbf{x}) \in [0, 1]. \quad (4)$$

Vectorizing $\Phi > 0$ over the coastal SAR footprint yields 201 polygons totalling 2.86 km² (Fig. 3). This area is a *lower bound* on flooding that occurred, not a flood map: it records only what was inundated at the instants the satellite happened to look.

C. CONDITIONING FACTORS AND COLLINEARITY SCREENING

Twelve candidate factors were computed (Fig. 4). TWI and the stream power index follow the standard definitions [20]:

$$\text{TWI} = \ln\left(\frac{a}{\tan \beta}\right), \quad \text{SPI} = \ln(a \tan \beta), \quad (5)$$

with a the specific catchment area (m²) and β the local slope. HAND is the elevation of a pixel above the nearest hydrologically connected drainage cell [21, 22].

TABLE II: Data sources. Every layer is public; none is hand-authored or simulated.

Source	Product	Resolution	Role
Sentinel-1 [13]	GRD IW, VV+VH, rel. orbit 63	10 m	Multi-temporal flood-frequency inventory
JRC GSW [26]	Water occurrence	30 m	Permanent-water exclusion mask
SRTM [45]	Digital elevation model	30 m	Elevation, slope, aspect, curvature
MERIT Hydro [46]	HAND, upstream area	90 m	HAND, TWI, distance to river, drainage density
CHIRPS [47]	Daily precipitation	5 km	Mean annual rainfall climatology
ESA WorldCover [48]	Land cover	10 m	Ablation only (excluded from primary model)
Sentinel-2	Surface reflectance	10 m	NDVI; ablation only
Open-Meteo / ERA5 [49]	Hourly precipitation	point	Live and archived antecedent rainfall
WorldTides	Tide height	point	Acquired for the live dashboard; <i>not</i> used in $T(t)$ (§V-J)
OpenStreetMap	Drivable road network	vector	Routing graph [50]
Natural Earth	Administrative outlines	10 m scale	Cartography only

GRD: ground range detected; IW: interferometric wide swath; VV/VH: co-polarized and cross-polarized channels; JRC GSW: Joint Research Centre Global Surface Water

Collinearity was screened by the variance inflation factor (VIF),

$$\text{VIF}_j = \frac{1}{1 - R_j^2}, \quad (6)$$

where R_j^2 is the coefficient of determination of factor j regressed on all others. As Table III shows, TWI and SPI are mutually collinear by construction, both being functions of α and $\tan \beta$; with both present they score 16.22 and 14.50. Removing SPI reduces TWI to 1.49 and leaves every retained factor below 3.5, comfortably under the conventional threshold of 10.

D. SAMPLING WITH TERRAIN-MATCHED NEGATIVES

Training points are drawn by stratified sampling, 1500 per class, under two constraints. Positives require $\Phi(\mathbf{x}) \geq 0.05$; negatives require $\Phi(\mathbf{x}) = 0$ and a distance of at least 120 m from any flood pixel, which removes the ambiguous halo around observed flooding. Both classes are confined to a floodable-lowland domain,

$$\mathcal{D} = \{\mathbf{x} : z(\mathbf{x}) \leq 280 \text{ m} \wedge h(\mathbf{x}) \leq 80 \text{ m}\}, \quad (7)$$

with z elevation and h HAND, both thresholds set at the 90th percentile of the flood-pixel distribution. Without $\mathcal{D}$ the negative class would be drawn largely from forested uplands, and the classifier could separate the classes on elevation alone without learning anything flood-specific. This elevation- and HAND-matching of negatives is, in our assessment, the single decision that most affects whether the reported scores mean anything.

E. A REMOTE-SENSING CONFOUND, AND ITS REMOVAL

The first model trained on all twelve factors reached a test area under the receiver operating characteristic curve (ROC-AUC) of 0.9878. SHAP attribution, however, ranked land cover and NDVI as the two most important predictors, ahead

of elevation and HAND (Fig. 7c). Inspecting the sample explained why: within the floodable-lowland domain, non-flood pixels are overwhelmingly closed-canopy vegetation. C-band SAR cannot observe standing water beneath a closed canopy [27, 28]. A pixel under forest is therefore labeled “not observed flooded” for a sensing reason, not a hydrological one, and a model given land cover can achieve high accuracy by learning *where SAR can see the ground*, not where water collects.

This is circularity, not signal, and it would have transferred a spurious model into an operational routing system. We therefore define the primary feature set as terrain and rainfall only, nine factors, and retain the twelve-factor model solely as a sensitivity analysis. Section V-A quantifies the cost of this decision: 0.9878 to 0.9625 in test ROC-AUC, with the SHAP ranking reverting to elevation, HAND, distance to river, rainfall and TWI, an ordering consistent with the physical literature.

F. CLASSIFIER, CALIBRATION AND EXPLANATION

Random forest [51] and XGBoost [52] were trained on a 70/30 stratified split, with 5-fold cross-validation on the training portion. XGBoost was selected on test ROC-AUC.

Raw tree-ensemble scores are ranking outputs, not probabilities [53], so we calibrate with isotonic regression [54] fitted by cross-validation on the training split alone, and evaluate with the Brier score [55]. Because the training sample is class-balanced while the landscape is not, the calibrated score is tied to a 0.5 prior; Section V-B quantifies this and applies the standard prior correction [56, 57]:

$$\text{BS} = \frac{1}{n} \sum_{i=1}^n (\hat{p}_i - y_i)^2, \quad (8)$$

with $\hat{p}_i$ the predicted probability and $y_i \in \{0, 1\}$ the label. Isotonic regression finds the monotone g minimizing $\sum_i (g(s_i) - y_i)^2$ over the raw scores s_i , so calibration cannot

TABLE III: Conditioning factors. VIF is reported both with SPI present (11 numeric candidates) and for the retained primary set. SPI is dropped because it and TWI are algebraically coupled through specific catchment area and slope; removing one resolves both.

Factor	Definition / units	VIF (11)	VIF (9)	Source
Elevation	Height above sea level (m)	3.87	3.46	SRTM
Slope	Terrain gradient (°)	3.72	1.54	SRTM
Aspect	Slope azimuth (°)	1.04	1.01	SRTM
Curvature	Laplacian of elevation (−)	1.10	1.10	SRTM
HAND	Height above nearest drainage (m)	2.91	2.85	MERIT Hydro
TWI	$\ln(a/\tan \beta)$ (−)	16.22	1.49	MERIT Hydro
Distance to river	Euclidean distance to channel (m)	2.71	2.71	MERIT Hydro
Drainage density	Channel fraction in 1 km neighborhood (−)	2.21	2.19	MERIT Hydro
Annual rainfall	2020–2024 mean (mm)	2.16	2.14	CHIRPS
<i>Dropped:</i>				
SPI	$\ln(a \tan \beta)$ (−)	14.50	–	VIF > 10
NDVI	Normalized difference vegetation index	1.52	–	Confound (Sec. IV-E)
Land cover	ESA WorldCover class	–	–	Confound (Sec. IV-E)

reorder predictions and therefore cannot alter AUC by construction.

Feature attribution uses SHAP values [10], the Shapley decomposition of a prediction over feature subsets:

$$\phi_j = \sum_{S \subseteq \mathcal{F} \setminus \{j\}} \frac{|\mathcal{S}|! (|\mathcal{F}| - |\mathcal{S}| - 1)!}{|\mathcal{F}|!} [f(\mathcal{S} \cup \{j\}) - f(\mathcal{S})], \quad (9)$$

with $\mathcal{F}$ the feature set; we report $|\phi_j|$ over the test split.

G. THE VALIDATION LADDER

Three successively stricter protocols are reported.

(i) Random split. Standard stratified 70/30 with 5-fold CV. Comparable to published practice, and optimistic for the reasons below.

(ii) Spatial-block cross-validation. Samples are assigned to 0.05° (≈ 5.5 km) blocks and whole blocks are held out [15], so no test point shares a neighborhood with a training point. The block size exceeds the terrain autocorrelation range in this setting.

(iii) Geographic transfer. The strictest test. The model is trained *only* on the 2589 labeled points outside the coastal SAR footprint plus a 0.01° buffer, then predicts the 21,365 valid pixels inside that footprint and is scored against the SAR polygons observed there. The 411 labeled points inside are withheld entirely. This is a strict within-district coastal hold-out, not evidence of cross-district transfer.

For (iii) we report ROC-AUC, hit rate (probability of detection), false-alarm ratio and critical success index:

$$\begin{aligned} \text{POD} &= \frac{\text{TP}}{\text{TP} + \text{FN}}, & \text{FAR} &= \frac{\text{FP}}{\text{TP} + \text{FP}}, \\ \text{CSI} &= \frac{\text{TP}}{\text{TP} + \text{FP} + \text{FN}}, \end{aligned} \quad (10)$$

together with the success-rate curve [58], the cumulative fraction of observed flooding captured as the landscape is ranked by descending susceptibility. Introduced for landslide-hazard validation and standard in susceptibility mapping since, the

success-rate curve is the appropriate headline metric here, because FAR and CSI are severely distorted by the fact that SAR-observed flooding covers only 0.95% of the window and absence of observation is not evidence of dryness.

H. DYNAMIC TRIGGER

The trigger maps antecedent rainfall to $[0, 1]$ through three terms referenced to India Meteorological Department daily-rainfall categories, without any parameter fitted to the events it is later tested on:

$$\begin{aligned} T(t) = \min \Big(& \underbrace{0.5 \frac{\min(P_{24}, \theta_{\text{VH}})}{\theta_{\text{VH}}}}_{\text{daily intensity}} + \underbrace{0.3 \frac{\min(P_{72}, \theta_{\text{sat}})}{\theta_{\text{sat}}}}_{\text{catchment saturation}} \\ & + \underbrace{0.2 \frac{\min(P_6, \theta_{\text{H}})}{\theta_{\text{H}}}}_{\text{short burst}}, 1 \Big), \end{aligned} \quad (11)$$

where P_6, P_{24}, P_{72} are trailing rainfall totals (mm) over 6, 24 and 72 h; $\theta_{\text{H}} = 64.5$ mm (IMD Heavy) and $\theta_{\text{VH}} = 115.6$ mm (IMD Very Heavy) are published category thresholds; and $\theta_{\text{sat}} = 200$ mm is a saturation reference. We note explicitly that while θ_{H} and θ_{VH} are published IMD category boundaries, the three weights and θ_{sat} are chosen by us: θ_{sat} is a round value on the order of the wet-season soil-moisture storage of these catchments, not a measured constant. Section V-D sweeps all four. The three-term structure encodes the physical premise that a saturated catchment floods on less additional rain. Equation (11) is monotone non-decreasing in each P by construction, a property we verify numerically rather than assert (Section V-D). The outer $\min(\cdot, 1)$ is redundant at the operating weights, since they sum to one and each term is individually capped, but is retained because Section V-D evaluates weightings for which it binds.

For live district operation, Eq. (11) is evaluated at Mangaluru, Moodbidri, Bantwal, Puttur and Sullia and the conservative maximum is applied district-wide. This improves monitoring coverage but does not remove the spatially uniform-trigger assumption. Tide is displayed separately and




FIGURE 2. Pipeline. The two-stage split into a static SAR-trained surface and a dynamic live trigger (shaded path through $S(x)$ and $T(t)$) is a direct consequence of the twelve-day single-orbit revisit established in Section I, not a modeling preference.

is not an input to $T(t)$.

I. FLOOD-AWARE GRAPH AND ROUTING

The exact district drivable network is extracted from OpenStreetMap (25,407 nodes, 59,501 directed edges) and each edge is sampled at its midpoint against the 120 m, nine-factor susceptibility raster produced by the serialized XGBoost model. Raster no-data cells are filled by nearest valid values solely for road sampling; no road midpoint fell outside the raster in this run. Edge cost under trigger T is

$$c(e) = \begin{cases} t(e)(1 + \kappa R(e)), & R(e) < \tau, \\ \infty, & R(e) \geq \tau, \end{cases} \quad (12)$$

$$R(e) = S(e) \cdot T,$$

with $t(e)$ free-flow travel time (s), κ a dimensionless penalty weight and τ the impassability threshold. The operating point is $\kappa = 8$, $\tau = 0.60$; Section V-F sweeps both.

Dijkstra [36] and A^* [35] are implemented explicitly, not called from a graph library, specifically so the set of nodes each expands is observable; that set is the object of comparison. The A^* heuristic is a straight-line travel-time lower bound,

$$h(u, g) = \frac{d_{\text{hav}}(u, g)}{v_{\text{max}}}, \quad v_{\text{max}} = 22 \text{ m s}^{-1}, \quad (13)$$

SAR flood inventory: 201 polygons, 2.9 km²




FIGURE 3. The multi-temporal SAR flood inventory over the coastal footprint: 201 polygons, 2.86 km², on a HAND backdrop with the channel network and estuary. Observed flooding concentrates in the low-HAND corridors flanking the Gurupura and Nethravathi and along the tidal margin.

with d_{hav} the haversine distance. Admissibility requires

$$h(u, g) \leq h^*(u, g) \quad \forall u, \quad (14)$$

where h^* is true optimal remaining cost. Since no edge in the graph permits travel faster than v_{max} and the haversine distance lower-bounds path length, (14) holds; we additionally verify it numerically against the graph's own edge weights before reporting any result, and cross-check both implementations against *networkx* as a third reference.

J. QUANTUM DISPATCH: FORMULATION

A 59,501-edge routing problem is far beyond exact quantum simulation, so the quantum layer is applied only where the problem is genuinely small: assigning n ambulances to n incidents.

It is worth being explicit about what this layer is and is not for, because the classical and quantum components of this paper serve different time horizons. The graph layer of Section IV-I is the operational component: it runs on commodity hardware, on a real network, today. Linear assignment as posed here is not a bottleneck in that pipeline. It is solvable exactly in $O(n^3)$ by the Hungarian algorithm [39], and at the fleet sizes a district of this size operates, the classical solver is exact and effectively instantaneous. No quantum method improves on that, and none is claimed to.

The reason to state the assignment as a quadratic unconstrained binary optimization problem at all is that the operational dispatch problem does not stay linear. The realistic extensions, capacitated vehicles, time windows, sequential re-assignment as the flood evolves, and multiple incidents per vehicle, are NP-hard, and QUBO or Ising form is the standard route by which such problems are posed for annealing and gate-model solvers [41, 59, 60]. Those formulations cannot be checked against an exact optimum at useful sizes, which




FIGURE 4. The nine conditioning factors of the primary model over the Mangaluru window. The channel network is visible as low-slope, high-TWI, low-HAND corridors; the coarse CHIRPS grid is apparent in the rainfall panel and is a known resolution limit of that product.

is precisely why the encoding, the penalty construction and the parameter-transfer schedule should be validated first on the case where ground truth is computable. That is the role of this section: it verifies the machinery, not the advantage. Establishing that the QUBO ground state coincides with the Hungarian optimum, and characterising how the approximation ratio behaves with circuit depth, is what would make a later comparison on NP-hard dispatch variants, or on physical hardware, interpretable rather than anecdotal.

Why $n = 3$ when the operational evaluation is district-wide: The one-hot encoding below needs n^2 qubits. Exact state-vector simulation and brute-force ground-state verification grow exponentially, whereas at $n = 3$ the state space is only

512 amplitudes and every check remains exact. We therefore solve a district-spread 3×3 instance: Wenlock District Hospital, Puttur General Hospital and Sullia Taluk Hospital are assigned to scenarios at Mulki, Belthangady and Subrahmanya. The cost matrix C comes from the same flood-aware routing layer used throughout, so the quantum stage consumes the paper's own risk surface rather than a synthetic instance. The reduction is a constraint of exact simulation and verification, not of the formulation: Equations (15)–(17) are stated for general n .

With binary $x_{ai} = 1$ iff ambulance a serves incident i , and risk-aware travel times C_{ai} from the classical layer, the

QUBO is

$$\min_{x \in \{0,1\}^{n^2}} \sum_{a,i} C_{ai} x_{ai} + P \sum_a \left(\sum_i x_{ai} - 1 \right)^2 + P \sum_i \left(\sum_a x_{ai} - 1 \right)^2, \quad (15)$$

with penalty $P = 2 \max_{a,i} C_{ai} + 1$, which is sufficient to make every constraint-violating state higher in energy than any feasible one. Substituting $x_j = (1 - z_j)/2$, $z_j \in \{-1, +1\}$ gives the Ising cost Hamiltonian [41]

$$H_C = \sum_{j < k} J_{jk} \sigma_j^z \sigma_k^z + \sum_j h_j \sigma_j^z + \text{const.} \quad (16)$$

QAOA [40] prepares

$$|\psi_p(\gamma, \beta)\rangle = \prod_{\ell=1}^p e^{-i\beta_\ell H_B} e^{-i\gamma_\ell H_C} |+\rangle^{\otimes n^2}, \quad H_B = \sum_j \sigma_j^x, \quad (17)$$

minimizing $\langle \psi_p | H_C | \psi_p \rangle$ over (γ, β) by constrained optimization by linear approximation (COBYLA) [61], a derivative-free trust-region method appropriate here because the expectation is evaluated exactly but has no cheap analytic gradient. Depth $p + 1$ is warm-started from depth p by the INTERP rule [42], linearly interpolating the optimized schedule onto a finer grid. Quality is reported as the approximation ratio

$$r = \frac{E_{\max} - \langle H_C \rangle}{E_{\max} - E_{\min}} \in [0, 1], \quad (18)$$

with $E_{\min}, E_{\max}$ from exact enumeration of all 2^{n^2} states. For $n = 3$ this is 9 qubits and 512 states, so the ground state is verified by brute force and against the Hungarian algorithm [39] on the raw cost matrix. Simulation is exact statevector; no physical quantum hardware is used.

Because COBYLA is a local optimizer on a non-convex landscape, a single run at depth p is one draw from a distribution, not the performance at that depth. We therefore repeat every depth over five random seeds and report mean and standard deviation.

K. IMPLEMENTATION

The pipeline is Python 3.13 with pinned dependencies: earthengine-api and geemap for server-side processing, scikit-learn [62], xgboost [52] and shap [10] for modeling, osmnx [50] and networkx for network acquisition and reference validation, geopandas/rasterio/shapely for geometry, numpy [63] and scipy [64] for the statevector simulation and the Hungarian reference solver, and matplotlib [65] for all figures. Artifacts are built by a dependency-aware Makefile, so a full rebuild regenerates every number and figure reported here from source data. A suite of 84 automated tests encodes the methodological invariants as executable assertions, including class balance and domain confinement of the sample, VIF < 10 for all retained factors,

monotonicity of $T(t)$, heuristic admissibility, agreement of Dijkstra, A* and networkx on every benchmarked path, and equality of the QUBO ground state with the Hungarian optimum. Every experiment reported here ran on one Apple M3 Pro workstation with 18 GB RAM. Apart from the server-side raster work that Earth Engine performs on its own infrastructure, no additional compute was needed.

V. RESULTS AND DISCUSSION

A. FEATURE-SET ABLATION: THE PRICE OF REMOVING CIRCULARITY

Table IV reports the ablation ladder. The twelve-factor model is the strongest by every metric, and is the model we discard. Its SHAP ranking places land cover first and NDVI second; once both are removed the ranking becomes elevation, HAND, distance to river, annual rainfall, TWI, which is the ordering the flood-susceptibility literature reports on independent study regions [8, 9, 11]. Dropping SPI on collinearity grounds costs essentially nothing (0.9878 to 0.9872), confirming it carried no independent information. Dropping NDVI and land cover costs 0.025 AUC and 0.038 F_1 .

We regard the 0.9625 figure as the honest one and the 0.9878 as an artifact of sensor physics. A reviewer comparing this work against papers reporting 0.98–0.99 on similar data should be aware that the difference may be a methodological choice, not a modeling deficiency.

A smaller version of the same problem, in a factor we retained: Having built the argument above, it would be inconsistent not to apply the same scepticism to the surviving features, and one does not fully survive it. CHIRPS has a 0.05° (≈ 5.5 km) grid, so the coastal footprint contains only 17 distinct rainfall values across 292,125 pixels, yet SHAP ranks `rain_annual` fourth. A variable that coarse cannot be resolving fine-scale hydrology. Measuring what it is actually doing: its correlation with distance from the coast is 0.73 and with elevation 0.33, so within this window it functions substantially as a smooth coast-to-inland gradient, that is, a proxy for the orographic ramp, not a local rainfall signal.

It is not, however, empty. Leave-one-out retraining shows removing it costs 0.0063 AUC, the second largest of any single factor after elevation (0.0173), so it carries information the terrain variables do not. Our reading is that at district scale the CHIRPS gradient encodes a real orographic precipitation contrast that matters for susceptibility, while at city scale it degenerates into a positional covariate. We retain it, because the model is trained district-wide, and flag it here so the SHAP ranking is not over-interpreted as evidence of fine-scale rainfall control. This is a weaker version of the NDVI problem: a variable earning attribution partly for a reason other than the one its name implies.

Model selection between random forest and XGBoost on the primary set is close: random forest reaches accuracy 0.9033 and test ROC-AUC 0.9618 (0.9554 ± 0.0110 under 5-fold cross-validation), XGBoost 0.9111 and 0.9625

TABLE IV: Feature-set ablation on the held-out test split ($n = 900$). ``12F`` includes SPI, NDVI and land cover; ``11F`` drops SPI on VIF grounds; ``9F`` is the primary terrain-and-rainfall model. All models are XGBoost with identical hyperparameters and split.

Set	Acc.	Prec.	Rec.	F_1	ROC-AUC	CV AUC	SHAP top-3
12F	0.9456	0.9368	0.9556	0.9461	0.9878	0.9839 ± 0.0092	lulc, ndvi, elevation
11F	0.9478	0.9390	0.9578	0.9483	0.9872	0.9833 ± 0.0098	lulc, ndvi, elevation
9F	0.9111	0.9167	0.9044	0.9105	0.9625	0.9585 ± 0.0154	elevation, hand, dist_river




FIGURE 5. SHAP attribution for the primary model: elevation and HAND dominate, followed by channel proximity, rainfall climatology and wetness.

(0.9534 ± 0.0086). XGBoost is selected on test ROC-AUC and used throughout. Its SHAP attribution (Fig. 5) ranks elevation and HAND first, followed by channel proximity, rainfall climatology and wetness, an ordering consistent with the hydrological reading of the terms and free of the land-cover dominance that characterized the discarded 12-factor model.

B. PROBABILITY CALIBRATION, AND WHAT THE SCORE DOES NOT MEAN

Isotonic calibration improved the Brier score from 0.0697 to 0.0683 and, being a monotone transform, left ranking performance unchanged (ROC-AUC 0.9648). The reliability diagram (Fig. 6) tracks the diagonal across the probability range.

It would be easy, and wrong, to conclude from this that $S(\mathbf{x})$ can be read as a probability of flooding. The training sample is stratified 1500/1500, so its implied base rate is exactly 0.5 by construction, and isotonic regression fitted on that split calibrates to a 0.5 prior. The deployment prevalence, measured as the SAR-observed flooded fraction of the coastal holdout, is 0.0095 (Table VII). The calibrated score is therefore well-behaved with respect to a reference population in which half the landscape floods, which is not the landscape.

The consequence is measurable and, importantly, it partly explains a result we initially attributed elsewhere. At a 0.5 cut the balanced-prior model flags 40.76% of the coastal holdout as positive. Applying the standard prior correction [56, 57], for a model trained at prior π_{tr} and deployed at π_{dep} ,

$$p' = \frac{p\omega}{p\omega + (1-p)}, \quad \omega = \frac{\pi_{dep}/(1-\pi_{dep})}{\pi_{tr}/(1-\pi_{tr})} = 0.009591, \quad (19)$$

drops that to 5.7524% (Table V). Because (19) is strictly in-

Reliability diagram (Brier 0.0697 $\rightarrow$ 0.0683)




FIGURE 6. Reliability diagram after isotonic calibration. Predicted probabilities track observed frequencies across the range.

creasing in p , ROC-AUC is invariant at 0.9648, which is the point: the ranking was never in question, only the interpretation of its absolute level.

Section V-C defends the high false-alarm ratio on the grounds that absence of SAR observation is not evidence of dryness. That argument stands, but it is not the whole explanation, and we correct it here: a substantial part of the 40.76% positive fraction is simply the balanced-prior artifact. We therefore report $S(\mathbf{x})$ throughout as a *relative exposure ranking*, which is what the routing layer consumes and what the success-rate curve evaluates, and we do not recommend thresholding it at 0.5 as though it were a posterior. An operational deployment wanting absolute probabilities should apply (19) with a locally estimated prevalence, bearing in mind that the SAR-derived π_{dep} used here is itself a lower bound.

C. SPATIAL VALIDATION: HOW MUCH SKILL IS REAL

Table VI and Fig. 7 report the validation ladder, and the pattern across it is the most methodologically interesting result in this paper.

Under spatial-block cross-validation the primary model falls from 0.9585 to 0.9474, a drop of 1.16%. The reduction is modest, indicating that the model does not substantially exploit spatial autocorrelation: withholding geography costs it approximately one point of AUC. Notably the twelve-factor model drops far *less* (0.30%), which is not evidence that it generalizes better. Land cover is so strongly predictive of the SAR label that it swamps the spatial structure,

TABLE V: Effect of the training prior. The classifier is trained on a balanced sample; the landscape is not balanced. The prior-corrected posterior changes what a score means, not how scores rank.

Quantity	Balanced prior	Prevalence-corrected
Implied base rate	0.500	0.0095
Brier score (balanced test split)	0.0683	—
ROC-AUC	0.9648	0.9648 (invariant)
Coastal holdout above a 0.5 cut	40.76%	5.7524%

so the block-CV control has little left to remove. A spatial-CV score taken in isolation would therefore have ranked the confounded model as the more robust one, which is a caution worth recording: spatial validation controls for leakage through geography, not for leakage through a sensor artifact.

Under the strictest test, coastal geographic transfer, AUC falls to 0.8594 (Table VII). POD is 0.9507, and the success-rate curve (Fig. 8b) shows the top 10% of the landscape by susceptibility contains 49.26% of all SAR-observed flooding and the top 20% contains 75.37%, lifts of $4.93\times$ and $3.77\times$ over a random ranking. This is the operationally defensible discrimination result. Fig. 9 shows the surfaces side by side; drainage-following structure persists in the withheld coastal footprint.

The false-alarm ratio of 0.9778 requires care, not concealment. Observed flood pixels are 0.95% of the footprint, so at a 0.5 threshold the majority of predicted-positive pixels are necessarily unobserved-flooded. More importantly, under a twelve-day revisit a pixel outside the polygons means “not seen flooded”, never “dry”. The FAR is therefore an upper bound on the true false alarm rate and should not be read as a 97% error rate; the AUC and the success-rate curve, both of which use the full ranking rather than a threshold, are the interpretable summaries.

D. TRIGGER VALIDATION AND ITS DISCRIMINATIVE LIMIT

Table VIII publishes the complete input record behind every trigger value, which we consider necessary: a trigger whose inputs are not shown cannot be checked. Equation (11) is verified monotone under uniform rainfall scaling ($T = 0.000, 0.063, 0.125, 0.251, 0.502, 0.924, 1.000$ for scale factors 0 to 8).

The ordering test passes. Both documented floods score above an ordinary wet day ($T = 0.133$) and a dry baseline ($T = 0.000$), with no parameter fitted to them. Notably the two events arrive at high T by different routes: May 2025 through daily intensity (87.2 mm/24 h, IMD Heavy) and August 2024 through saturation (216.9 mm/72 h) despite a merely Moderate 24-hour total. A single-threshold rainfall alarm keyed to 24-hour totals would have ranked August 2024 below an ordinary heavy day; the three-term structure is what recovers it.

The adversarial case is where the result is weaker, and we report it as measured. A deliberately chosen heavy-rain day with no reported flooding (15 July 2025; 66.8 mm/24 h, IMD Heavy) scores $T = 0.551$, only 0.011 below the August 2024 flood. The trigger separates flood from non-flood conditions

strongly but separates flood from *heavy-rain-without-flood* barely at all. Fig. 10b shows why, and also suggests the remedy: the discriminating variable is the 72-hour saturation total (216.9 and 205.6 mm for the two floods versus 116.8 mm for the non-flood day), while the burst term actively works against separation, being highest of all on the non-flood day (28.0 mm/6 h). Re-weighting toward saturation would widen the margin, but doing so using these events would convert a validated trigger into a fitted one, so we report the limitation rather than tune it away.

Sensitivity of the trigger's own constants

Equation (11) contains four author-chosen constants: the three weights and θ_{sat} . Only θ_H and θ_{VH} are externally published. Having swept the routing constants, it would be inconsistent not to sweep these, particularly since they govern the paper's weakest result. Table IX reports 231 weight triples on the simplex at 0.05 resolution crossed with five values of θ_{sat} , evaluated on the same five scenarios.

Two findings follow. First, the *easy* discrimination is invariant to the weighting: floods rank above the ordinary wet day and the dry baseline in all 1155 settings, so that result is a property of the input data rather than of the weights. Second, the *hard* discrimination is governed almost entirely by one weight. The adversarial margin rises monotonically from -0.200 at zero saturation weight to $+0.416$ at unit saturation weight (Fig. 10c), and is positive in only 374 of 1155 settings. The operating point, at saturation weight 0.3, sits essentially on the zero crossing, which is why its margin is a marginal $+0.011$.

This converts the qualitative diagnosis of the previous subsection into a measurement: the 72-hour antecedent term carries the flood-versus-heavy-rain signal, and the operating weights under-use it. We deliberately do *not* adopt the optimum. Choosing weights to maximize separation on the same five events the trigger is validated against would convert a validated component into a fitted one and forfeit the property this paper claims for it. Re-weighting requires an independent event set, which we identify as future work.

Why tidal stage is the physically indicated fourth term

Re-weighting the existing three terms is unlikely to be the right resolution, because all three are rainfall statistics and the quantity they omit is not rainfall. In a tidally influenced estuary the water level at a given discharge is set jointly by the discharge and by the seaward boundary condition. A high tide raises that boundary, flattens the hydraulic gradient at the mouth, and reduces the discharge the channel can con-



TABLE VI: Random versus spatial-block 5-fold cross-validation. Blocks are 0.05° squares; whole blocks are held out.

Feature set	Random 5-fold AUC	Spatial-block AUC	Drop	Drop (%)
12F	0.9839 ± 0.0092	0.9810 ± 0.0052	0.0030	0.30
11F	0.9833 ± 0.0098	0.9805 ± 0.0059	0.0028	0.29
9F	0.9585 ± 0.0154	0.9474 ± 0.0039	0.0111	1.16

TABLE VII: Coastal geographic-transfer validation. The model is trained on the 2589 labeled points outside the SAR footprint plus a 0.01° buffer and scored against SAR polygons inside it; 411 inside points are withheld entirely. Contingency metrics use a 0.5 threshold.

Quantity	Value	Quantity	Value
Training points (outside)	2589	ROC-AUC	0.8594
Points withheld (inside)	411	Hit rate (POD)	0.9507
Valid prediction pixels	21,365	False-alarm ratio [†]	0.9778
Observed-flood pixels	203 (0.95%)	Critical success index [†]	0.0221
TP / FN	193 / 10	Flooding in top 10% area	49.26%
TN / FP	12,647 / 8515	Flooding in top 20% area	75.37%

[†]Upper and lower bounds respectively: SAR observes flooding only on its twelve-day revisit, so a pixel outside the polygons is “not observed flooded”, not “dry”.

**FIGURE 7.** (a) Feature-set ablation. (b) Random versus spatial-block cross-validation; the confounded 12F model drops least, because land cover dominates the spatial signal. (c) SHAP top-6 for the 9-factor and 12-factor models side by side, showing land cover and NDVI displacing terrain in the confounded model.**FIGURE 8.** Coastal geographic holdout. (a) ROC against SAR-observed flooding for a model trained entirely outside the evaluation footprint and its buffer. (b) Success-rate curve: the highest-susceptibility 10% of the landscape contains 49.26% of observed flooding.




FIGURE 9. (a) The calibrated susceptibility surface $S(x)$ over the coastal footprint with SAR-observed flooding outlined. (b) The geographic-transfer prediction from a model trained outside this footprint and its buffer. The two surfaces are visually similar, which is the qualitative counterpart of the transfer AUC of 0.8594.

TABLE VIII: Complete trigger inputs, from the ERA5 archive via Open-Meteo at Mangaluru (12.87°N, 74.84°E). Weighted terms are the three components of (11). No parameter is fitted to these events.

Scenario	Class	Rainfall (mm)			IMD (24 h)	Weighted terms			T
		6 h	24 h	72 h		daily	sat.	burst	
May 2025 Mangaluru urban	flood	7.6	87.2	205.6	Heavy	0.377	0.300	0.024	0.701
Aug 2024 Nethravathi river	flood	10.1	53.4	216.9	Moderate	0.231	0.300	0.031	0.562
15 Jul 2025 heavy rain	no flood	28.0	66.8	116.8	Heavy	0.289	0.175	0.087	0.551
6 Sep 2025 ordinary wet	wet	3.2	11.0	50.0	Light	0.048	0.075	0.010	0.133
16 Feb 2025 dry baseline	dry	0.0	0.0	0.0	No rain	0.000	0.000	0.000	0.000




FIGURE 10. (a) Decomposition of $T(t)$ into its three weighted terms. (b) The measured rainfall inputs. The two documented floods exceed 200 mm of 72-hour antecedent rainfall; the heavy-rain non-flood day reaches only 116.8 mm despite the largest 6-hour burst in the set, identifying catchment saturation as the discriminating variable. (c) Mean adversarial margin across the weight simplex as a function of the saturation weight, at $\theta_{\text{sat}} = 200$ mm. The operating point sits almost exactly at the zero crossing.

vey, so an identical rainfall sequence produces a higher inland stage when its runoff peak coincides with high water.

This is the standard fluvial-tidal compound mechanism, and it is the mechanism the present formulation cannot represent:

$T(t)$ in (11) is a function of P_6, P_{24}, P_{72} alone, so any two days with matching rainfall receive identical T regardless of tidal phase.

Two properties of the required extension follow from that mechanism. First, the tidal term should enter multiplicatively rather than additively. An additive fourth term would let a high tide alone raise T on a dry day, which is physically wrong; the effect is conditional on there being discharge to impound. A form such as

$$T'(t) = T(t) \cdot (1 + \lambda \eta^+(t)), \quad \eta^+(t) = \max\left(0, \frac{\eta(t) - \eta_0}{\eta_{\max} - \eta_0}\right), \quad (20)$$

with $\eta(t)$ the observed stage, η_0 a threshold above which backwater becomes material, and λ a single coupling constant, has the property that it vanishes when either rainfall or tidal anomaly is absent and grows only when both are present. Second, η should be measured rather than predicted. Harmonic tide predictions of the kind the live system currently retrieves give the astronomical component only; the meteorological residual, the storm surge and wind setup that accompany the same weather systems that produce the rainfall, is exactly the part correlated with flood events, and it is present only in gauge observations.

We state this as a hypothesis with a defined test, not as a result. The test is whether η^+ evaluated at the runoff peak separates the two documented floods from the 15 July heavy-rain day at a margin materially larger than the +0.011 obtained from rainfall alone. We have not run it, because it requires gauge records for the Gurupura and Nethravathi mouths that we could not obtain for the study period, and λ and η_0 would themselves need calibration on events independent of the five used here. Reporting the mechanism and the test without the result is the honest position: the +0.011 margin is a real limit on the current trigger, and we can name the variable most likely to explain it without claiming to have shown that it does.

E. FLOOD-AWARE ROUTING

At the standardized May 2025 setting ($T = 0.70$, $\kappa = 8$, $\tau = 0.60$), 4019 directed edges are blocked. Static susceptibility exceeds 0.5 on 9886 of 59,501 edges (16.61%; Fig. 11). For the representative Wenlock Hospital–Kulur pair, the flood-aware route increases travel time from 9.5 min to 16.7 min while reducing mean exposure from 0.193 to 0.073 (62.2%) and maximum exposure from 0.690 to 0.545 (Fig. 12). Both routes still depend on a small set of river crossings, showing that software cannot substitute for physical redundancy.

The district design contains 247 hospital–scenario pairs. Of these, 192 remain reachable under the flood-aware policy and 55 are severed. On the 192 survivors, mean exposure falls from 0.1712 (risk-blind) to 0.0655 (risk-aware), a 61.7% reduction, with a mean detour of 63.075 min. This conditional average must be read together with the severed count; averaging only the routes that remain feasible otherwise hides the largest operational consequence.

Robustness to scenario selection: The 13 named scenario points are author-selected. Repeating the analysis for 200 road nodes sampled uniformly at random yields 2,865 routed and 935 severed pairs out of 3,800. Among routed pairs, mean exposure falls from 0.1828 to 0.0696 (61.9%) with a mean 71.177 min detour. The exposure reduction is therefore similar under randomized destinations, while the large severed count confirms that district-scale connectivity, rather than favorable site choice, dominates the operational result.

F. SENSITIVITY OF THE ROUTING CONSTANTS

Fig. 13 and Table X report the full 7×7 sweep of (κ, τ) , re-computing all 247 hospital–scenario pairs in each cell.

Correcting for survivorship: At $\tau \leq 0.50$, 87 of the 247 pairs are severed, so averaging “whatever routed” would discard the hard cases. All comparative statistics below and in Fig. 13(a)–(b) are therefore computed over the 160 pairs routable in every cell; panel (c) reports severed counts over all 247. The risk-blind baseline on that common subset is mean exposure 0.1619 with mean travel time 55.194 min.

The penalty weight remains material: At $\tau = 0.60$ over the common subset, increasing κ from 2 to 16 lowers mean exposure from 0.0646 to 0.0552 while increasing mean detour from 56.902 min to 73.471 min. The eight-cell neighborhood around the operating point has a 22.76% exposure spread. Thus the earlier claim that κ becomes immaterial is not supported at district scale; $\kappa = 8$ is a documented operating choice within a continuing safety–delay trade-off.

The impassability threshold is a different kind of parameter: Holding $\kappa = 8$, severed pairs are 87 at $\tau \leq 0.50$, 55 at $\tau \in \{0.55, 0.60\}$, 31 at $\tau = 0.65$, and 0 at $\tau \geq 0.70$. Passability turns on a small number of critical edges whose risk values cluster tightly. Both parameters matter: κ controls a continuous exposure–detour trade-off, while τ creates discrete changes in district connectivity. The threshold should be calibrated against observed road closures and set by policy on acceptable risk rather than inherited as convention.

G. DIJKSTRA VERSUS A* : A MULTI-QUERY BENCHMARK

Table XI and Fig. 14 report all 192 pairs reachable under the standardized May 2025 setting. The two searches recover the same path and cost for every pair. A* settles fewer nodes on all 192: mean exploration falls from 10,830 to 9344.1, a 13.7% mean reduction, with a median per-pair reduction of 10.03% and a range of 0.33–80.81%.

The node saving does not translate into faster execution in this transparent Python implementation. Mean wall time rises from 29.715 ms for Dijkstra to 36.033 ms for A*, so A* is 21.3% slower on average and faster on only 53 of 192 pairs. Haversine-heuristic evaluation costs more than the saved expansions on many queries. These timings are

TABLE IX: Sensitivity of the trigger to its four unpublished constants, over 231 weight triples $\times$ 5 values of θ_{sat} . ``Adversarial margin" is  $\min_{\text{floods}} T$  minus  $T$  for the heavy-rain non-flood day; ``ordinary margin" is against the ordinary wet day. The optimum is reported for information and is not adopted.

Setting	Weights (daily/sat./burst), θ_{sat}	Adversarial	Ordinary
Operating point (adopted)	0.50 / 0.30 / 0.20, 200 mm	+0.011	+0.430
Best over the sweep (not adopted)	0.00 / 1.00 / 0.00, 200 mm	+0.416	+0.610
Worst over the sweep	0.00 / 0.00 / 1.00, 100 mm	−0.316	+0.062
Settings with correct flood > ordinary > dry ordering		1155 / 1155	
Settings with a positive adversarial margin		374 / 1155	
Settings with an adversarial margin above 0.05		277 / 1155	




FIGURE 11. (a) The routing graph colored by per-edge susceptibility. (b) The 9886 flood-prone edges ($S > 0.5$), 16.61% of the network. Statistics are computed on the full district graph of 59,501 directed edges.

TABLE X: Sensitivity sweep, selected sections. Exposure and detour are computed over the 160 pairs routable in every cell of the sweep, so columns are comparable; the severed count is over all 247. Risk-blind baseline on that common subset: mean exposure 0.1619 and mean travel time 55.194 min.

$\tau = 0.60$ fixed			$\kappa = 8$ fixed			
κ	Exposure	Detour (min)	τ	Exposure	Detour (min)	Severed
2	0.0646	56.902	0.45	0.0532	69.050	87
4	0.0646	57.190	0.50	0.0550	68.384	87
6	0.0636	58.838	0.55	0.0563	67.362	55
8	0.0624	60.944	0.60	0.0624	60.944	55
10	0.0588	63.685	0.65	0.0757	24.969	31
12	0.0568	70.314	0.70	0.0923	13.434	0
16	0.0552	73.471	0.75	0.0923	13.434	0

implementation-bound, but they are the measurements produced by the tested reference code.

H. QUANTUM DISPATCH: A FEASIBILITY CHECK

The QUBO ground state matches the Hungarian optimum exactly in both scenarios: 124.7229 min for dry conditions and 250.8706 min under the May 2025 trigger. Table XII reports the depth study over five seeds.

Mean approximation ratio is non-monotone with depth in

both scenarios. The dry case peaks at $p = 3$ (0.9107 ± 0.0096) and then falls, while the flood case is highest at $p = 1$ (0.9004 ± 0.0199) and varies thereafter. Optimum-sampling probability is also non-monotone, ranging from 0.0274 to 0.0978 for dry and 0.0615 to 0.1461 for flood. COBYLA is a local optimizer on a non-convex landscape and the INTERP warm start does not guarantee improvement, so circuit depth cannot be presented as a monotone benefit here.

The optimal assignment is identical under dry and flood




FIGURE 12. Risk-blind (orange) versus flood-aware (green) route from Wenlock District Hospital to Kulur at $T = 0.70$, over the susceptibility-colored network. The flood-aware route takes an inland detour, trading 7.2 min for a 62.2% reduction in mean exposure.

(a)-(b) over the 160 pairs routable in every cell (risk-blind baseline exposure 0.162); operating point $\kappa=8$, $\tau=0.6$ (red box)




FIGURE 13. Sensitivity of the routing outcome to the penalty weight κ and impassability threshold τ at $T = 0.70$. (a) Mean route exposure and (b) mean detour, both over the 160 pairs routable in every cell of the sweep so that cells are compared like with like; (c) severed pairs, over all 247. Vertical banding shows the connectivity discontinuities controlled by τ ; the red box marks the operating point.

conditions: Wenlock–Mulki, Puttur–Belthangady and Sullia–Subrahmanya. We had expected re-allocation and did not observe it: at this fleet size and geometry, flooding changes travel times enough to alter route *shape* (Section V-E) but not enough to alter the assignment’s relative ordering. The finding is reported as observed.

We state plainly that at $n = 3$ the Hungarian algorithm is exact and effectively instantaneous, while the QAOA sim-

ulation takes seconds and returns the optimum only probabilistically. This layer is a feasibility demonstration of the QUBO/QAOA formulation on a real risk-aware cost matrix, not a claim of quantum advantage. Its value is that the encoding, the constraint penalties and the verification path are in place should problem sizes and hardware make the question live.

TABLE XI: Dijkstra versus A* across 192 reachable hospital–scenario pairs at $T = 0.70$.

Metric	Dijkstra	A*
Mean nodes explored	10,830	9344.1
Median per-pair node reduction	10.03%	
Range of per-pair reduction	0.33% – 80.81%	
Pairs where A* explores fewer	192 / 192	
Mean wall time	29.715 ms	36.033 ms
Pairs where A* is faster	53 / 192	
Identical optimal path and cost	192 / 192	




FIGURE 14. (a) Per-pair search cost; points below the diagonal indicate A* explored fewer nodes. (b) Distribution of per-pair node-exploration reduction, median 10.03%.

TABLE XII: QAOA depth study, 9 qubits, exact statevector, mean $\pm$ s.d. over 5 seeds with INTERP warm starting. "Lift" is mean optimum-sampling probability relative to the uniform baseline $1/2^9$.

p	Scenario	Params	Approx. ratio	Best	P(optimum)	Lift
1	Dry	2	0.8759 ± 0.0078	0.8899	0.0279 ± 0.0414	14.3×
2	Dry	4	0.8896 ± 0.0114	0.9007	0.0353 ± 0.0300	18.1×
3	Dry	6	0.9107 ± 0.0096	0.9192	0.0978 ± 0.0845	50.1×
4	Dry	8	0.8953 ± 0.0078	0.9062	0.0701 ± 0.0857	35.9×
5	Dry	10	0.8953 ± 0.0077	0.9097	0.0274 ± 0.0414	14.0×
1	Flood	2	0.9004 ± 0.0199	0.9211	0.0615 ± 0.0765	31.5×
2	Flood	4	0.8954 ± 0.0122	0.9114	0.0719 ± 0.0705	36.8×
3	Flood	6	0.8882 ± 0.0083	0.8968	0.0699 ± 0.0617	35.8×
4	Flood	8	0.8998 ± 0.0119	0.9197	0.1461 ± 0.1051	74.8×
5	Flood	10	0.8957 ± 0.0109	0.9104	0.0795 ± 0.0794	40.7×

I. COMPARISON WITH PRIOR WORK

Table XIII positions the reported performance against representative susceptibility studies. Direct comparison of headline AUCs across study regions is weak evidence, since inventory density, class balance and sampling design differ, but the table makes the intended point: our random-split figure sits inside the published range, and the figures that follow from stricter validation, which most of these studies do not report, sit below it. We would encourage the comparison to be read column-wise on protocol rather than row-wise on score.

J. LIMITATIONS

- Susceptibility is not depth.** $S(x)$ and $R(x, t)$ are relative exposure indices in $[0, 1]$, calibrated to a balanced prior rather than to landscape prevalence (Section V-B), so τ is a heuristic on that index and not a statement about whether a vehicle can ford a segment. A depth regression in the manner of Sadhwani and Eldho [30], who obtain depth from SAR and elevation over this same coastal belt, is the clearest route to a defensible passability criterion.
- Tide is not in the trigger.** Equation (11) is rainfall-only. This is the most conspicuous omission in the paper, because the Nethravathi and Gurupura discharge into a

TABLE XIII: Reported flood-susceptibility performance in representative prior work and this study. “Protocol” indicates the strictest validation reported.

Study	Method	Reported AUC	Protocol
Tehrany et al. [6]	SVM + frequency ratio	0.88	Random split
Khosravi et al. [7]	Decision-tree family	0.85–0.92	Random split
Arabameri et al. [8]	Random forest	0.93	Random split
Saha et al. [9]	Hyperpipes + SVR ensemble	0.94	Random split
Aydin and Iban [11]	Boosting ensembles + SHAP	> 0.95	Random split
Sadhwani and Eldho [30]	Dual ML, SAR + elevation	—	Random split
This study (9F, primary)	XGBoost + SHAP	0.9625	Random split
This study (9F)	XGBoost + SHAP	0.9474	Spatial-block CV
This study (9F)	XGBoost + SHAP	0.8594	Coastal holdout

Values are as reported by each study on its own region and inventory and are not directly comparable; the column that carries the argument is **Protocol**. Our discarded 12-factor model scores 0.9878 on a random split, which would top this table, and is excluded because Section V-A shows that score is a sensor artifact.

shared estuary where tidal stage can impound river flow, so a high tide coinciding with peak discharge is precisely the compound condition the framework should capture. Tide height is acquired for the live system but does not enter any reported T value. Adding a fourth term keyed to tidal stage at the flood peak is the single highest-value extension to this work.

3. **The trigger is spatially uniform.** The live system samples five sites and conservatively aggregates their maximum trigger, but then applies that scalar uniformly, so scenarios differ in intensity but never in spatial pattern. Equation (1) therefore assumes the relative spatial pattern of risk is invariant to event type, which cannot hold for events that are fluvially versus tidally driven.
4. **Trigger discrimination is weak at the top of the range.** As Section V-D shows, the margin between a documented flood and a heavy-rain non-flood day is 0.011. The trigger is usable for distinguishing flood conditions from normal ones and should not currently be used as a flood/no-flood alarm.
5. **The inventory is a lower bound.** With a twelve-day revisit, absence of SAR-observed flooding at a location is never evidence of safety. Every contingency metric computed against it inherits this asymmetry.
6. **Validation is single-district.** All results are from Dakshina Kannada. The geographic-transfer test withholds a coastal footprint inside the district; it does not cross a climatic or geomorphic regime.
7. **Road risk is a raster midpoint approximation.** The operational surface is 120 m resolution and samples each edge at its midpoint; it is neither a hydraulic depth simulation nor a guarantee of conditions along the full segment.
8. **Hospital origins are not a live capability roster.** The 19 curated OSM locations were audited on 2026-09-11 for routing provenance, not for ambulance availability, emergency specialty, beds or referrals.

9. **The quantum layer demonstrates feasibility only.** At $n = 3$ the classical optimum is exact and instant; no advantage is claimed, and the simulation is noiseless, so it does not characterize noisy intermediate-scale quantum (NISQ) behavior.

10. **Wall-clock routing results are implementation-bound.** The pure-Python router was chosen for search transparency; its timing should not be read as a statement about A^* in production systems.

VI. CONCLUSION

This paper set out to build a flood-risk and emergency-response framework for a coastal Indian district whose satellite record is thinner than the methods literature usually assumes, and to be explicit about what that record can and cannot support. Five findings follow.

1. **A large part of the apparent skill of flood-susceptibility models in forested coastal catchments can be a sensor artifact.** Including land cover and NDVI raised test ROC-AUC from 0.9625 to 0.9878, but SHAP attribution showed the gain came from those two variables, and the mechanism is that C-band SAR cannot see water under closed canopy in a landscape whose non-flooded lowland is predominantly forest. The model was learning detectability, not hazard. Removing the confound costs 0.025 AUC and is, we argue, mandatory.
2. **Spatial validation controls for geography, not for sensor circularity, and can invert the ranking of models.** Under spatial-block CV the primary model lost 1.16% AUC while the confounded model lost only 0.30%, because land cover swamps the spatial signal. Under strict coastal geographic holdout, training entirely outside the evaluation footprint and its buffer, the primary model achieves ROC-AUC 0.8594, with the top 10% of the landscape by susceptibility containing 49.26% of SAR-observed flooding and the top 20% containing 75.37%.

3. **A trigger built only from published rainfall thresholds orders documented events correctly, but discriminates poorly at the top of its range.** With no parameter fitted to them, the May 2025 and August 2024 floods score $T = 0.701$ and 0.562 against 0.133 for an ordinary wet day and 0.000 for a dry baseline; the two events reach high T by different mechanisms, intensity and saturation respectively, which a 24-hour threshold alarm would miss. But a heavy-rain day with no reported flooding scores 0.551 , a margin of only 0.011 , and the diagnostic points at the 72-hour saturation term as the variable carrying the real signal.
4. **District-wide routing exposes a sharp safety–access trade-off.** Of 247 pairs, 192 remain reachable and 55 are severed at the operating point. On the survivors, mean exposure falls from 0.1712 to 0.0655 (61.7%) at a mean 63.075 min detour. The full 7×7 sweep shows that κ remains material, while τ creates discrete connectivity changes: 87, 55, 31 and 0 severed pairs across successive threshold regimes.
5. **Search-transparent benchmarking exposes trade-offs that single examples hide.** A* returned the identical optimal path to Dijkstra on all 192 reachable pairs and explored 13.7% fewer nodes on average, yet was 21.3% slower and faster on only 53 queries. The district-spread 3×3 QUBO reproduces the Hungarian optimum, but QAOA performance is non-monotone with depth and offers no demonstrated advantage.

The broader argument is methodological. Each stricter protocol reduced the headline number, 0.9878 to 0.9625 to 0.9474 to 0.8594 , and each removed a specific, nameable way the earlier figure was optimistic. For a practitioner assessing whether to trust such a map, that sequence is more informative than its first term alone.

Future work follows directly from the limitations. A calibrated depth regression would replace the heuristic impassability threshold with a hydraulic criterion and make τ a physical quantity instead of a policy one. A spatially resolved rainfall field would make the trigger genuinely spatiotemporal and address the strongest assumption in (1). Re-weighting the trigger toward catchment saturation should widen the flood versus heavy-rain margin, but requires an independent event set to remain a validation rather than a fit. Extending the geographic-transfer test across districts with different geomorphology would establish whether the model transfers regionally or only locally. Finally, executing the QAOA circuits on physical NISQ hardware would characterize the noise behavior that exact simulation cannot.


REFERENCES

- [1] V. D. Chandrashekar and A. Shetty, "Trends in extreme rainfall over ecologically sensitive Western Ghats and coastal regions of Karnataka: an observational assessment," *Arabian Journal of Geosciences*, vol. 11, no. 12, Art. no. 327, Jun. 2018, doi: 10.1007/s12517-018-3700-6.
- [2] S. Chandana and M. C. Chandan, "Hydrological Implications of LULC Changes in the Netravathi River Basin: Insights from SWAT Modeling and Machine Learning Techniques," in *2023 IEEE India Geoscience and Remote Sensing Symposium (InGARSS)*, pp. 1–4, Dec. 2023, doi: 10.1109/InGARSS59135.2023.10490324.
- [3] S. Babar and H. Ramesh, "Streamflow Response to Land Use–Land Cover Change over the Nethravathi River Basin, India," *Journal of Hydrologic Engineering*, vol. 20, no. 10, Art. no. 05015002, Oct. 2015, doi: 10.1061/(ASCE)HE.1943-5584.0001177.
- [4] N. Naik, K. Chandrasekaran, V. M. Sundaram, and P. Panneer, "Dual attention guided deep encoder-decoder network for change analysis in land use/land cover for Dakshina Kannada District, Karnataka, India," *Environmental Earth Sciences*, vol. 82, no. 1, Art. no. 33, Dec. 2022, doi: 10.1007/s12665-022-10713-1.
- [5] R. Nirmala, "Flood susceptibility assessment of Nethravathi River basin, Karnataka, through GIS-AHP technique. In *Sustainability Solutions*, pages 643–661. Springer, 2026, doi: 10.1007/978-3-032-11300-9_27.
- [6] M. S. Tehrani, B. Pradhan, and M. N. Jebur, "Flood susceptibility mapping using a novel ensemble weights-of-evidence and support vector machine models in GIS," *Journal of Hydrology*, vol. 512, pp. 332–343, May 2014, doi: 10.1016/j.jhydrol.2014.03.008.
- [7] K. Khosravi et al., "A comparative assessment of decision trees algorithms for flash flood susceptibility modeling at Haraz watershed, northern Iran," *Science of The Total Environment*, vol. 627, pp. 744–755, Jun. 2018, doi: 10.1016/j.scitotenv.2018.01.266.
- [8] A. Arabameri, K. Rezaei, A. Cerdà, C. Conoscenti, and Z. Kalantari, "A comparison of statistical methods and multi-criteria decision making to map flood hazard susceptibility in Northern Iran," *Science of The Total Environment*, vol. 660, pp. 443–458, Apr. 2019, doi: 10.1016/j.scitotenv.2019.01.021.
- [9] A. Saha et al., "Flood Susceptibility Assessment Using Novel Ensemble of Hyperpipes and Support Vector Regression Algorithms," *Water*, vol. 13, no. 2, p. 241, Jan. 2021, doi: 10.3390/w13020241.
- [10] S. M. Lundberg, S.-I. Lee, "A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems 30*, 2017.
- [11] H. E. Aydin and M. C. Iban, "Predicting and analyzing flood susceptibility using boosting-based ensemble machine learning algorithms with SHapley Additive exPlanations," *Natural Hazards*, vol. 116, no. 3, pp. 2957–2991, Dec. 2022, doi: 10.1007/s11069-022-05793-y.
- [12] S. T. Seydi, Y. Kanani-Sadat, M. Hasanlou, R. Sahraei, J. Chanut-sot, and M. Amani, "Comparison of Machine Learning Algorithms for Flood Susceptibility Mapping," *Remote Sensing*, vol. 15, no. 1, p. 192, Dec. 2022, doi: 10.3390/rs15010192.
- [13] R. Torres et al., "GMES Sentinel-1 mission," *Remote Sensing of Environment*, vol. 120, pp. 9–24, May 2012, doi: 10.1016/j.rse.2011.05.028.
- [14] M. Arora, S. Sahoo, C. M. Bhatt, P. K. Litoria, and B. Pateriya, "Rapid flood inundation mapping and impact assessment using Sentinel-1 SAR data over Ghaggar River basin of Punjab, India," *Journal of Earth System Science*, vol. 132, no. 4, Art. no. 183, Nov. 2023, doi: 10.1007/s12040-023-02199-7.
- [15] D. R. Roberts et al., "Cross-validation strategies for data with temporal, spatial, hierarchical, or phylogenetic structure," *Ecography*, vol. 40, no. 8, pp. 913–929, Mar. 2017, doi: 10.1111/ecog.02881.
- [16] P. Ploton et al., "Spatial validation reveals poor predictive performance of large-scale ecological mapping models," *Nature Communications*, vol. 11, no. 1, Art. no. 4540, Sep. 2020, doi: 10.1038/s41467-020-18321-y.
- [17] "Review for "Predicting into unknown space? Estimating the area of applicability of spatial prediction models",", Sep. 2020, doi: 10.1111/2041-210x.13650/v1/review2.
- [18] Y. Wang, Z. Fang, H. Hong, and L. Peng, "Flood susceptibility mapping using convolutional neural network frameworks," *Journal of Hydrology*, vol. 582, p. 124482, Mar. 2020, doi: 10.1016/j.jhydrol.2019.124482.
- [19] M. Avand, H. Moradi, and M. R. Iasboyee, "Using machine learning models, remote sensing, and GIS to investigate the effects of changing climates and land uses on flood probability," *Journal of Hydrology*, vol. 595, p. 125663, Apr. 2021, doi: 10.1016/j.jhydrol.2020.125663.


- [20] K. J. Beven, M. J. Kirkby, A physically based, variable contributing area model of basin hydrology. *Hydrological Sciences Bulletin*, 24(1):43–69, 1979.
- [21] C. D. Rennó et al., “HAND, a new terrain descriptor using SRTM-DEM: Mapping terra-firme rainforest environments in Amazonia,” *Remote Sensing of Environment*, vol. 112, no. 9, pp. 3469–3481, Sep. 2008, doi: 10.1016/j.rse.2008.03.018.
- [22] A. D. Nobre et al., “Height Above the Nearest Drainage – a hydrologically relevant new terrain model,” *Journal of Hydrology*, vol. 404, no. 1–2, pp. 13–29, Jun. 2011, doi: 10.1016/j.jhydrol.2011.03.051.
- [23] S. Martinis, A. Twele, and S. Voigt, “Towards operational near real-time flood detection using a split-based automatic thresholding procedure on high resolution TerraSAR-X data,” *Natural Hazards and Earth System Sciences*, vol. 9, no. 2, pp. 303–314, Mar. 2009, doi: 10.5194/nhess-9-303-2009.
- [24] A. Twele, W. Cao, S. Plank, and S. Martinis, “Sentinel-1-based flood mapping: a fully automated processing chain,” *International Journal of Remote Sensing*, vol. 37, no. 13, pp. 2990–3004, Jun. 2016, doi: 10.1080/01431161.2016.1192304.
- [25] X. Shen, D. Wang, K. Mao, E. Anagnostou, and Y. Hong, “Inundation Extent Mapping by Synthetic Aperture Radar: A Review,” *Remote Sensing*, vol. 11, no. 7, p. 879, Apr. 2019, doi: 10.3390/rs11070879.
- [26] J. Pekel, A. Cottam, N. Gorelick, and A. S. Belward, “High-resolution mapping of global surface water and its long-term changes,” *Nature*, vol. 540, no. 7633, pp. 418–422, Dec. 2016, doi: 10.1038/nature20584.
- [27] P. A. Townsend, “Relationships between forest structure and the detection of flood inundation in forested wetlands using C-band SAR,” *International Journal of Remote Sensing*, vol. 23, no. 3, pp. 443–460, Jan. 2002, doi: 10.1080/01431160010014738.
- [28] V. Tsyganskaya, S. Martinis, P. Marzahn, and R. Ludwig, “SAR-based detection of flooded vegetation – a review of characteristics and approaches,” *International Journal of Remote Sensing*, vol. 39, no. 8, pp. 2255–2293, Jan. 2018, doi: 10.1080/01431161.2017.1420938.
- [29] G. Donchyts, F. Baart, H. Winsemius, N. Gorelick, J. Kwadijk, and N. v. d. Giesen, “Earth’s surface water change over the past 30 years,” *Nature Climate Change*, vol. 6, no. 9, pp. 810–813, Aug. 2016, doi: 10.1038/nclimate3111.
- [30] K. Sadhwani and T. I. Eldho, “Machine learning-driven flood hazard assessment: integrating SAR and elevation data for inundation mapping and depth estimation,” *Natural Hazards*, vol. 122, no. 5, Art. no. 177, Feb. 2026, doi: 10.1007/s11069-025-07949-y.
- [31] S. Thasma, Y.-H. Ho, and A. Mostafavi, “Beyond Backscatter: AlphaEarth Land-Cover Priors for Rapid SAR Flood Segmentation Across Foundation Backbones,” arXiv:2606.29134, 2026.
- [32] F. H. Siddiqui, M. M. Islam, and S. B. Noor, “A Temporal Planning Approach for Intelligent Flood Response,” arXiv:2608.20510, 2026.
- [33] N. Zhang, H. Alipour, Multi-scale analysis of the impact of flooding on the urban road network. *Transportation Research Part D*, 77:279–293, 2019.
- [34] J. Yin, D. Yu, Z. Yin, M. Liu, Q. He, Evaluating the impact and risk of pluvial flash flood on intra-urban road network. *Journal of Hydrology*, 537:138–145, 2016.
- [35] P. Hart, N. Nilsson, and B. Raphael, “A Formal Basis for the Heuristic Determination of Minimum Cost Paths,” *IEEE Transactions on Systems Science and Cybernetics*, vol. 4, no. 2, pp. 100–107, 1968, doi: 10.1109/tssc.1968.300136.
- [36] E. W. Dijkstra, “A Note on Two Problems in Connexion with Graphs,” *Edsger Wybe Dijkstra*, pp. 287–290, Jul. 2022, doi: 10.1145/3544585.3544600.
- [37] R. Church and C. R. Velle, “THE MAXIMAL COVERING LOCATION PROBLEM,” *Papers in Regional Science*, vol. 32, no. 1, pp. 101–118, Jan. 1974, doi: 10.1111/j.1435-5597.1974.tb00902.x.
- [38] L. Brotcorne, G. Laporte, and F. Semet, “Ambulance location and relocation models,” *European Journal of Operational Research*, vol. 147, no. 3, pp. 451–463, Jun. 2003, doi: 10.1016/S0377-2217(02)00364-8.
- [39] H. W. Kuhn, “The Hungarian method for the assignment problem,” *Naval Research Logistics (NRL)*, vol. 52, no. 1, pp. 7–21, Oct. 2004, doi: 10.1002/nav.20053.
- [40] E. Farhi, J. Goldstone, and S. Gutmann, “A quantum approximate optimization algorithm,” 2014, arXiv:1411.4028.
- [41] A. Lucas, “Ising formulations of many NP problems,” *Frontiers in Physics*, vol. 2, 2014, doi: 10.3389/fphy.2014.00005.
- [42] L. Zhou, S. Wang, S. Choi, H. Pichler, and M. D. Lukin, “Quantum Approximate Optimization Algorithm: Performance, Mechanism, and Implementation on Near-Term Devices,” *Physical Review X*, vol. 10, no. 2, Art. no. 021067, Jun. 2020, doi: 10.1103/physrevx.10.021067.
- [43] J. R. McClean, S. Boixo, V. N. Smelyanskiy, R. Babbush, and H. Neven, “Barren plateaus in quantum neural network training landscapes,” *Nature Communications*, vol. 9, no. 1, Art. no. 4812, Nov. 2018, doi: 10.1038/s41467-018-07090-4.
- [44] N. Gorelick, M. Hancher, M. Dixon, S. Ilyushchenko, D. Thau, and R. Moore, “Google Earth Engine: Planetary-scale geospatial analysis for everyone,” *Remote Sensing of Environment*, vol. 202, pp. 18–27, Dec. 2017, doi: 10.1016/j.rse.2017.06.031.
- [45] “Shuttle Radar Topography Mission (SRTM),” *Fact Sheet*, 2009, doi: 10.3133/fs20093087.
- [46] D. Yamazaki, D. Ikeshima, J. Sosa, P. D. Bates, G. H. Allen, and T. M. Pavelsky, “MERIT Hydro: A High-Resolution Global Hydrography Map Based on Latest Topography Dataset,” *Water Resources Research*, vol. 55, no. 6, pp. 5053–5073, Jun. 2019, doi: 10.1029/2019wr024873.
- [47] C. Funk et al., “The climate hazards infrared precipitation with stations—a new environmental record for monitoring extremes,” *Scientific Data*, vol. 2, no. 1, Art. no. 150066, Dec. 2015, doi: 10.1038/sdata.2015.66.
- [48] D. Zanaga et al. ESA WorldCover 10 m 2020 v100. Zenodo, 2021, doi: 10.5281/zenodo.5571936.
- [49] H. Hersbach et al., “The ERA5 global reanalysis,” *Quarterly Journal of the Royal Meteorological Society*, vol. 146, no. 730, pp. 1999–2049, Jun. 2020, doi: 10.1002/qj.3803.
- [50] G. Boeing, “OSMnx: New methods for acquiring, constructing, analyzing, and visualizing complex street networks,” *Computers, Environment and Urban Systems*, vol. 65, pp. 126–139, Sep. 2017, doi: 10.1016/j.compenvurbsys.2017.05.004.
- [51] “Random forests,” *Physics Subject Headings (PhySH)*, doi: 10.29172/7c2a6982-6d72-4cd8-bba6-2fccb06a7011.
- [52] T. Chen, C. Guestrin, XGBoost: A scalable tree boosting system. In *Proc. 22nd ACM SIGKDD*, pages 785–794, 2016.
- [53] A. Niculescu-Mizil and R. Caruana, “Predicting good probabilities with supervised learning,” in *Proceedings of the 22nd international conference on Machine learning - ICML ’05*, pp. 625–632, 2005, doi: 10.1145/1102351.1102430.
- [54] B. Zadrozny and C. Elkan, “Transforming classifier scores into accurate multiclass probability estimates,” in *Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining*, pp. 694–699, Jul. 2002, doi: 10.1145/775047.775151.
- [55] G. W. BRIER, “VERIFICATION OF FORECASTS EXPRESSED IN TERMS OF PROBABILITY,” *Monthly Weather Review*, vol. 78, no. 1, pp. 1–3, Jan. 1950, doi: 10.1175/1520-0493(1950)078<0001:vofeit>2.0.co;2.
- [56] C. Elkan, The foundations of cost-sensitive learning. In *Proc. 17th International Joint Conference on Artificial Intelligence (IJCAI)*, pages 973–978, 2001.
- [57] M. Saerens, P. Latinne, and C. Decaestecker, “Adjusting the Outputs of a Classifier to New a Priori Probabilities: A Simple Procedure,” *Neural Computation*, vol. 14, no. 1, pp. 21–41, Jan. 2002, doi: 10.1162/089976602753284446.
- [58] C. F. Chung and A. G. Fabbri, “Validation of Spatial Prediction Models for Landslide Hazard Mapping,” *Natural Hazards*, vol. 30, no. 3, pp. 451–472, Nov. 2003, doi: 10.1023/B:NHAZ.0000007172.62651.2b.
- [59] F. Glover, G. Kochenberger, and Y. Du, “Quantum Bridge Analytics I: a tutorial on formulating and using QUBO models,” *4OR*, vol. 17, no. 4, pp. 335–371, Dec. 2019, doi: 10.1007/s10288-019-00424-y.
- [60] D. Fitzek, T. Ghandriz, L. Laine, M. Granath, and A. Frisk Kockum, “Applying quantum approximate optimization to the heterogeneous vehicle routing problem,” *Scientific Reports*, vol. 14, Art. no. 25415, Oct. 2024, doi: 10.1038/s41598-024-76967-w.
- [61] M. J. D. Powell, A direct search optimization method that models the objective and constraint functions by linear interpolation. In *Advances in Optimization and Numerical Analysis*, pages 51–67. Springer, 1994, doi: 10.1007/978-94-015-8330-5_4.

- [62] F. Pedregosa, G. Varoquaux, A. Gramfort, et al. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- [63] C. R. Harris et al., “Array programming with NumPy,” *Nature*, vol. 585, no. 7825, pp. 357–362, Sep. 2020, doi: 10.1038/s41586-020-2649-2.
- [64] P. Virtanen et al., “SciPy 1.0: fundamental algorithms for scientific computing in Python,” *Nature Methods*, vol. 17, no. 3, pp. 261–272, Feb. 2020, doi: 10.1038/s41592-019-0686-2.
- [65] J. D. Hunter, “Matplotlib: A 2D Graphics Environment,” *Computing in Science & Engineering*, vol. 9, no. 3, pp. 90–95, 2007, doi: 10.1109/MCSE.2007.55.


AUTHOR BIOGRAPHIES




MUSTAFA BASTHIKODI [AUTHOR TO SUPPLY: degree(s) with field, awarding institution and year; current position or programme; research interests; IEEE membership grade if any. Delete these brackets.]




KRISH P. CHOWTA [AUTHOR TO SUPPLY: degree(s) with field, awarding institution and year; current position or programme; research interests; IEEE membership grade if any. Delete these brackets.]



ADVAITH SHAJEENDRAN [AUTHOR TO SUPPLY: degree(s) with field, awarding institution and year; current position or programme; research interests; IEEE membership grade if any. Delete these brackets.]



DHRITHI K. S. [AUTHOR TO SUPPLY: degree(s) with field, awarding institution and year; current position or programme; research interests; IEEE membership grade if any. Delete these brackets.]



GAMAN A. RAI [AUTHOR TO SUPPLY: degree(s) with field, awarding institution and year; current position or programme; research interests; IEEE membership grade if any. Delete these brackets.]

VOLUME 12, 2024